

PRAKHAR SHARMA

+91-8755887760 ◇ Jaipur, Rajasthan, India ◇ *Open to Bengaluru / Remote*

Gmail ◇ LinkedIn ◇ GitHub ◇ Portfolio

SUMMARY

AI Engineer designing, building, and deploying production **Generative AI solutions: RAG pipelines** on vector databases, **LLM-powered chat systems and AI agents**, and automated multi-step workflows shipped to real users. Hands-on with **LangChain, LlamaIndex, and HuggingFace Transformers**, including **fine-tuning and evaluating LLMs**. Experienced in authoring **reusable AI skills, system rules, guardrails, and prompt templates** that standardize LLM behavior in production. Strong Python (FastAPI) and Node.js, deploying containerized services via **Docker on AWS** and Azure with CI/CD.

SKILLS

LLM & GenAI	Large Language Model (LLM) Applications, Generative AI (GenAI), OpenAI API (GPT-4o), Google Gemini, LangChain, LlamaIndex, HuggingFace Transformers, LLM Fine-Tuning & Evaluation, AI Agents & Agentic Workflows, Function Calling / Tool Use, Structured Outputs (Pydantic), Prompt Engineering, Prompt Templates, Guardrails & Reusable AI Skills / System Instructions
Retrieval	Retrieval-Augmented Generation (RAG) Pipelines, Vector Databases (Pinecone), Vector Embeddings, Document Chunking & Ingestion, Semantic Search, Query Caching
Languages & Backend	Python, FastAPI, Node.js, Express.js, JavaScript / TypeScript, REST APIs (RESTful), Microservices Architecture, MongoDB (Mongoose), Redis
Cloud & MLOps	Docker, Docker Compose, Amazon Web Services — AWS (EC2, S3), Microsoft Azure (MSAL, Microsoft Graph API), CI/CD (Continuous Integration / Deployment), Git / GitHub, Automated Testing, BullMQ, PM2
Reliability	Dead-Letter Queues (DLQ), Transactional Outbox, Webhook & Send Idempotency, Distributed Schedulers (Mongo claim/lease), Pipeline Monitoring & Audit Logging, Role-Based Access Control (RBAC), OpenAPI / Swagger
Real-Time & Voice	Voice AI (Bolna), LiveKit, WebRTC, Socket.IO, Plivo, Twilio
Frontend	Next.js, React.js, Redux Toolkit, Tailwind CSS, Vitest

EXPERIENCE

AI Engineer

Nov 2025 – Present

The Odin, Jaipur, Rajasthan

- Designed, built, and deployed production **RAG pipelines** on **OpenAI and Pinecone**: automated **data ingestion**, chunking, and embedding into a company knowledge base; semantic search with a query-cache layer; and **LLM chat systems** grounded in retrieved company data, backed by a **post-LLM grounding validator** that corrects any figure disagreeing with the retrieval layer before it reaches the user.
- Created and versioned **reusable AI skills, system rules, guardrails, and workflow instructions** governing assistant behavior across chat, recruitment, and voice-agent workflows — including **prompt templates** and schema-constrained **structured outputs** that standardize LLM behavior and keep responses reliable and machine-parseable across internal and client-facing systems.
- Developed **autonomous AI voice agents** for candidate screening using **Bolna** and Python — webhook-driven call lifecycle handling, transcript capture, and **structured-output extraction** that writes screening results back into the ATS.
- Designed and shipped an **LLM mail-scheduling agent** (Religence, Aug 2026): chat turns natural language into a versioned **command contract** — the LLM never sends email or mutates data; a deterministic **Mongo**

claim/lease scheduler executes recurring sequences with per-run **idempotency keys**, Microsoft Graph correlation IDs, confirmation-gated activation, and classified retries.

- Engineered **production-ready AI pipelines** with **BullMQ/Redis** job queues, **dead-letter queues** for failed LLM jobs, the transactional outbox pattern, and webhook **idempotency guards** — with audit logging and continuous monitoring keeping async AI processing observable and safe to retry in production.
- Architected and shipped **Dharwin**, a multi-tenant applicant-tracking (ATS) / HRM SaaS platform — **99 data models and 77 REST modules** across **1,200+ commits** in six months (Node.js/Express/MongoDB + Next.js/TypeScript), integrating AI capabilities into existing systems via **REST APIs and microservices**: AI task breakdown, LiveKit video interviews, RBAC-scoped chat, and multi-provider email sync.
- Shipped **Religence**, a pharma lead-generation platform pairing a **Python/FastAPI** discovery service — web crawling with **OpenAI-based classification** and Pydantic-validated structured outputs — with a TypeScript/Next.js CRM, an in-CRM **Outlook inbox** (Microsoft Graph delta sync), and templated outreach.

PROJECTS

Dharwin — AI-Powered ATS/HRM Platform [GitHub] *Node.js · Express · MongoDB · Next.js · Pinecone · Bolna · Docker* Production SaaS platform unifying applicant tracking, HR operations, an LMS, and project management. Implemented a RAG knowledge base on Pinecone, conversational AI assistants with grounded-count validation, AI task breakdown, LiveKit video interviews, and Gmail + Microsoft Graph email sync. Includes **autonomous voice screening agents** (Bolna + OpenAI) that call candidates, handle the full webhook call lifecycle, and extract schema-constrained interview data — feeding transcripts and qualification status straight into the ATS. 99 models, 77 REST modules, 81 frontend tests.

Religence — Pharma Lead Discovery & CRM [GitHub] *Python · FastAPI · TypeScript · Next.js · OpenAI · Docker* End-to-end lead generation platform: a containerized Python/FastAPI crawler classifying prospect websites via OpenAI with Pydantic-validated structured outputs and Google Places enrichment, feeding a TypeScript CRM with verification queues, quotations, encrypted-token Outlook outreach (Microsoft Graph), and an **LLM mail-scheduling chatbot** — natural language to recurring email sequences via a versioned command contract executed by an idempotent lease-based scheduler (deployed Aug 2026).

Legal Document Generator [GitHub] *Flask · Gemini API · Pydantic* AI-powered legal document generation system using **Gemini** with structured outputs, function calling, and prompt orchestration to support multilingual document creation.

PROWPLUS — Project Management Portal [GitHub] *Node.js · Next.js 15 · npm Workspaces · shadcn/ui* Monorepo project-management portal that doubles as an embeddable client help-and-support portal (cross-site iframe sessions via httpOnly refresh cookies). Shared domain package (ticket stage-transition matrix, transactional email templates) consumed by both ends; validated ticket lifecycle with kanban and per-project analytics, dual-audience RBAC with internal-only comments and leak-hardened external client responses; 79 test files.

EDUCATION

B.Tech in Computer Science, AI & ML Specialization 2021 – 2025
GLA University, Mathura *Relevant coursework: Machine Learning (ML), Deep Learning, Natural Language Processing (NLP), Distributed Systems, Data Structures & Algorithms, Database Systems*

CERTIFICATIONS

- Machine Learning with Python** [Credly] Oct 2023
- AWS Educate: Introduction to Generative AI** [Credly] Jun 2025